

A New Algorithm for Calibrating Probit Models with Many Alternatives

Peter Cotton*

August 2026 — working draft, revised

Abstract

Multinomial probit choice probabilities are high-dimensional Gaussian integrals with no closed form; the standard tool, the GHK simulator, computes them one alternative at a time by smoothed importance sampling. For the factor-structured case $\Sigma = VV^\top + \text{diag}(D)$ we describe a reproducible approximation that computes *all* N probabilities together: conditionally on k Gaussian factors the alternatives are independent, and a lattice survival-product identity yields the entire share vector in $O(QNL)$ for Q factor nodes and L lattice points — Gauss–Hermite quadrature at low factor rank, fixed-design RQMC at higher rank. In tested $k = 2$ problems the maximum observed discrepancy from simulation references stayed below 10^{-3} from $N = 5$ to $N = 5000$ (21 seconds at $N = 5000$); a direct utility simulator at matched wall time reaches errors $1.6\text{--}6\times$ larger at these budgets and separates decisively below the 10^{-3} level, so the case for the method is error decay, reproducibility, resampling-free derivatives, and calibration — not raw forward speed alone. The share map is the gradient of the social-surplus function; its Jacobian is minus a weighted graph Laplacian, giving symmetric positive-definite reduced systems, $O(QNL)$ Jacobian–vector products, and a globally well-posed share inversion (existence, uniqueness modulo location) solved at $N = 1000$ in under a minute, with all N single-removal counterfactuals in $O(QN^2L)$ from one conditional field. The method requires the factor form: where Σ has no adequate low-rank-plus-diagonal approximation in the choice-relevant quotient space, simulation at exact Σ remains the tool.

1 Introduction

When a consumer chooses one product from a large assortment, a commuter one route from a network, or a platform ranks one item above thousands of others, the question a model must answer is: if an alternative disappears or changes, *where does its demand go?* The tools used for this, essentially universally, are the logit family — multinomial logit, nested logit, and above all random-coefficients (mixed) logit in the BLP tradition, the workhorse of demand estimation in empirical industrial organization, antitrust, and pricing [11, 4, 12]. They are universal for one reason: they are the ones that can be computed. Choice modeling did not start there. It

*Numbers in this paper are produced by committed, seeded scripts with correctness anchors run before any comparison (<https://github.com/microprediction/kinetics>, experiments 13–17; the exact commit is recorded in the repository README and an immutable tag is pinned at circulation). The algorithm ships in the **thurstone** package; package–benchmark agreement is itself a reported number. Benchmarks: Apple M4, 16 GB, Python 3.12.9, NumPy 2.4.6, SciPy 1.18.0; experiment 17 enforces single-threaded BLAS via **threadpoolctl** and reports median-of-3 timings for sub-second calls; long simulations are single realizations, disclosed as such. Monte Carlo chunk sizes derive from an explicit memory budget (1.5 GB); peak resident memory stays below 2 GB.

started with Thurstone’s 1927 model [1] — jointly Gaussian utilities, unrestricted correlation, hence unrestricted substitution — known in econometrics as multinomial probit (MNP). Its probabilities are $(N - 1)$ -dimensional Gaussian integrals with no closed form; the standard evaluator, the GHK simulator [2, 3, 4], is a separate smoothed importance-sampling run *per alternative* with $R^{-1/2}$ error in the draws. That cost is why probit survives mainly in small- N settings (transportation panels, a handful of alternatives) and effectively vanished from large- N practice: the field kept the computable models and gave up the model it started from. The concession is not free. Logit-family substitution is structurally restricted, and the whole apparatus of share inversion that demand estimation is built on exists on the logit side only. In the substitution study of Section 3, plain logit misallocates about 23% of the demand released by a deleted alternative where the correctly structured probit model misallocates 5%.

This paper’s point is that for the covariance structure applied work most often actually has — a factor model, $\Sigma = VV^\top + \text{diag}(D)$ — the lock is not real. We give one algorithm that solves *both directions* of the probit share map, plus its derivatives and counterfactuals:

1. **Forward pass** (utilities $\rightarrow$ shares): all N choice probabilities at once in $O(QNL)$ for Q factor nodes and L lattice points. All 5000 shares of an $N = 5000$ problem in 21 seconds, at accuracy near the 10^{-4} noise floor of the simulation reference; matching that with our per-alternative GHK implementation extrapolates to roughly forty hours.
2. **Calibration** (observed shares $\rightarrow$ utilities): the probit analogue of the logit-side inversions used throughout demand estimation [11, 12]. Globally well-posed (Theorem 1: existence and uniqueness modulo location) and solved for $N = 1000$ in under a minute.
3. **Derivatives**: Jacobian–vector products in $O(QNL)$ with an exact weighted-graph-Laplacian structure — symmetric positive-definite reduced systems, hence Krylov methods — with no resampling noise.
4. **Counterfactuals**: all N single-removal share vectors from one conditional field pass in $O(QN^2L)$.

The mechanism is a single trick used four ways: conditionally on the k factors the alternatives are independent, and a lattice survival-product identity computes every alternative’s integral in one shared $O(NL)$ pass — the field’s survival product is built once and each alternative’s leave-one-out integrand is recovered by division.

Scope, honestly. The method needs the factor form: where Σ has no adequate low-rank-plus-diagonal approximation in the choice-relevant quotient space, simulation at exact Σ remains the tool (Section 3). And at coarse accuracy targets ($\sim 10^{-3}$), plain direct utility simulation is a genuinely competitive forward baseline (Section 2); the decisive advantages are error decay below that level, reproducibility, resampling-free derivatives, and the inversion and counterfactual machinery. The one-pass identity itself comes from an unlikely literature — inferring ability from win probabilities in multi-entrant races [10] — which is perhaps why it has not been assembled for probit before, despite each ingredient being classical.

This paper concerns the factor-structured case

$$U_i = \mu_i + v_i^\top f + \sqrt{D_i} \varepsilon_i, \quad f \sim N(0, I_k), \quad \varepsilon_i \sim N(0, 1) \text{ i.i.d.}, \quad D_i > 0, \quad (1)$$

i.e. $\Sigma = VV^\top + \text{diag}(D)$. Factor dimension reduction for probit is itself classical — one-factor quadrature goes back to Butler and Moffitt [5], factor-analytic probit to Elrod and Keane [6]

— and GHK is accurate and well understood at moderate N . The contribution here is the all-alternatives assembly, its inverse transform, the derivative and counterfactual structure, and benchmark plus boundary evidence for all of it.

Proposition 1 (Conditional identity). *Under (1), with $g_i(\cdot | f)$ and $F_i(\cdot | f)$ the conditional density and CDF of U_i given f ,*

$$p_i = \mathbb{E}_f \left[\int g_i(x | f) \prod_{j \neq i} F_j(x | f) dx \right], \quad \sum_i p_i = 1, \quad (2)$$

and the product over $j \neq i$, for all i simultaneously, is obtained from $\log F_{\text{field}} = \sum_j \log F_j$ by subtraction.

Proof. Conditional on f the U_j are independent, so $\mathbb{P}(U_i \text{ maximal} | f) = \int g_i \prod_{j \neq i} F_j dx$; average over f . The events partition (ties are null for continuous laws), giving $\sum_i p_i = 1$. $\square$

Proposition 2 (Complexity). *With Q factor nodes and L lattice points, the forward share vector costs $O(QNL)$; the full single-removal ensemble $\{P(j \text{ chosen} | i \text{ removed})\}_{i,j}$ costs $O(QN^2L)$ arithmetic and $O(N^2)$ storage, reusing one conditional field construction.*

Proof. Per node: the field log-product costs $O(NL)$, each alternative's integral $O(L)$ after subtraction, giving $O(NL)$; the removal ensemble evaluates N survivor integrals for each of N removals, $O(N^2L)$, over the same cached field. $\square$

Proposition 3 (Identification and differentiability). *$p(\mu + c\mathbf{1}) = p(\mu)$, so $J = \partial p / \partial \mu$ has $\mathbf{1}$ in its null space and μ is identified only up to location. With B an $N \times (N - 1)$ basis of $\mathbf{1}^\perp$ and $\mu = B\eta$, share inversion and implicit differentiation are taken in reduced coordinates: $d\eta/d\theta = -(B^\top JB)^{-1} B^\top \partial p / \partial \theta$, provided $B^\top JB$ is nonsingular.*

Proof. Invariance holds because a common location shift does not change the argmax; differentiating $p(\mu + c\mathbf{1}) = p(\mu)$ in c gives $J\mathbf{1} = 0$. Differentiating $p(B\eta(\theta), \theta) = \hat{p}$ in θ and projecting with $B^\top$ gives the reduced formula under the stated regularity. $\square$

Theorem 1 (Weighted-Laplacian structure and global inversion). *For model (1) with $D_i > 0$, in max-wins coordinates, define for $i \neq j$*

$$w_{ij} = \mathbb{E}_f \int g_i(x | f) g_j(x | f) \prod_{\ell \neq i,j} F_\ell(x | f) dx > 0.$$

Then $J = \partial p / \partial \mu$ satisfies $J_{ij} = -w_{ij}$ and $J_{ii} = \sum_{j \neq i} w_{ij}$, i.e.

$$J = \sum_{i < j} w_{ij} (e_i - e_j)(e_i - e_j)^\top, \quad h^\top J h = \sum_{i < j} w_{ij} (h_i - h_j)^2,$$

a weighted graph Laplacian with everywhere-positive weights: J is symmetric, positive semidefinite, has $\mathbf{1}$ as its only null direction, and $B^\top JB \succ 0$ for every full-rank basis B of $\mathbf{1}^\perp$. Moreover $p = \nabla G$ for the social-surplus function $G(\mu) = \mathbb{E} \max_i \{\mu_i + \xi_i\}$, and for every strictly positive target $\hat{p}$ the inverse problem $p(\mu) = \hat{p}$ has exactly one solution on $\mathbf{1}^\perp$.

Proof. Differentiating (2) in μ_j and integrating by parts gives $J_{ij} = -w_{ij}$ for $i \neq j$; the diagonal follows from $J\mathbf{1} = 0$ (Proposition 3), and positivity of w_{ij} from positivity of Gaussian densities. The quadratic form is then immediate, and it vanishes only for constant h . For inversion, minimize $H(\mu) = G(\mu) - \hat{p}^\top \mu$ over $\mathbf{1}^\perp$: H is strictly convex there (its Hessian is the reduced J), and since the shocks have zero mean, $G(\mu) \geq \max_i \mu_i$, so

$$H(\mu) \geq \max_i \mu_i - \hat{p}^\top \mu = \sum_i \hat{p}_i (\max_j \mu_j - \mu_i) \geq \hat{p}_{\min} (\max_i \mu_i - \min_i \mu_i) \rightarrow \infty$$

as $\|\mu\| \rightarrow \infty$ on $\mathbf{1}^\perp$. The unique minimizer satisfies $\nabla G = \hat{p}$. (Every identity here is verified numerically in the repository: the w_{ij} representation to 6×10^{-9} against finite differences, symmetry to 4×10^{-12} , and positive definiteness of the reduced numerical Jacobian; the theorem statement was proposed by a referee.) Social-surplus convex duality for discrete choice is classical [15, 16]; the contribution here is not well-posedness itself but its scalable realization. $\square$

Proposition 4 (Jacobian–vector products in $O(QNL)$). *With hazards $\lambda_j = g_j/F_j$, $\Lambda = \sum_j \lambda_j$, and $A = \sum_j h_j \lambda_j$ (all conditional on f , all functions of x), the Jacobian of the exact max-wins map acts on a direction h as*

$$(Jh)_i = \mathbb{E}_f \int g_i \left(\prod_{j \neq i} F_j \right) (h_i \Lambda - A) dx,$$

where the field sums Λ and A are formed once per node and lattice point, so all N components cost $O(NL)$ per node. Translation invariance is visible ($h = \mathbf{1}$ gives $A = \Lambda$, hence $Jh = 0$), and the weighted-Laplacian form of Theorem 1 follows by expanding A . This enables Krylov solves of the reduced system in Proposition 3 without forming J densely at $O(QN^2L)$.

Proof. Differentiate (2): the μ_i -derivative of g_i contributes $\int g'_i \prod_{j \neq i} F_j$, which after integration by parts equals $\int g_i \sum_{j \neq i} \lambda_j F_j^{-1} \cdot (\prod_{\ell \neq i} F_\ell) g_j / \lambda_j$ -terms assembled as $\int g_i \prod_{j \neq i} F_j (\Lambda - \lambda_i)$; the μ_j -derivatives of the F_j factors contribute $-w_{ij} h_j$ summed to $-\int g_i \prod_{j \neq i} F_j (A - h_i \lambda_i)$. Combining gives the display. $\square$

Which map is differentiated, exactly. Three maps must be distinguished: the exact max-wins map $p^{\max}(\mu)$; the reflected min-wins map $p^{\min}(a)$ used by the implementation, related by $a = -\mu$ so that $D_\mu p^{\max}[h] = D_a p^{\min}[-h]$; and the implemented approximation $p_{Q,L} = \tilde{p}_{Q,L} / (\mathbf{1}^\top \tilde{p}_{Q,L})$, whose derivative carries a normalization correction: with $v = D\tilde{p}[h]$ and $s = \mathbf{1}^\top \tilde{p}$,

$$Dp_{Q,L}[h] = \frac{v - p_{Q,L} (\mathbf{1}^\top v)}{s}.$$

Since $\mathbf{1}^\top v = 0$ analytically, the correction is of the order of the quadrature defect; empirically the uncorrected formula matches central finite differences of the returned normalized map to 5×10^{-9} . One further convention: the adaptive lattice endpoints depend on μ , so derivative statements refer to the *fixed-design* map — the lattice envelope and factor nodes are frozen during any inversion or outer-derivative computation, and the inversion’s analytic own-location slopes are accordingly a Jacobi quasi-Newton preconditioner rather than the diagonal of the adaptive map’s Jacobian.

Tail-stable implementation. A referee exhibited a modest problem ($N = 8$, heterogeneous D) on which our published derivative code returned NaN/infinity: the density was floored before its logarithm while the exact log-survival stayed finite, producing spurious e^{+92} hazards. The implementation now works in the log domain throughout — $\log \Phi$ directly (never $1 - \Phi$, which underflows at $z \approx 8.3$), analytic $\log g$, hazards as $\exp(\log g - \log \Phi)$ (bounded, $\sim z$ in the tail) — and the failing case is a regression test. The same fix applies to the forward map, where deep-tail shares now resolve to genuine small positives rather than floored zeros.

The factors need not be Gaussian: (2) requires only conditional independence and an integrable factor law, so Student, mixture, or empirical factor distributions are admissible with appropriate nodes, and the base densities may be arbitrary and alternative-specific. One caution for asymmetric laws: the implementation works in the reflected (min-wins) race, and $-U$ has the same distribution as the reflected model only when the factor and base laws are symmetric; for asymmetric laws the reflected distributions must be used explicitly.

Numerical method, stated completely. The benchmarks use $L = 1501$ lattice points on the adaptive interval $[\min_{i,q} m_i(f_q) - 8 \max_i \sqrt{D_i}, \max_{i,q} m_i(f_q) + 8 \max_i \sqrt{D_i}]$, where $q = 1, \dots, Q$ indexes the retained factor nodes. Factor integration: the GHK benchmark (experiment 13) uses product Gauss–Hermite of order 15/11/9 for $k = 2/3/4$; the boundary experiment (experiment 14) uses order 15 per dimension for all $k \leq 4$, which after pruning weights below 10^{-7} of the maximum leaves $Q = 145, 1317$, and 10929 nodes for $k = 2, 3, 4$ — making the exponential growth of Q with rank explicit; scrambled-Sobol nodes (2^{13} points, fixed seed) are used for $k > 4$: reproducible fixed-seed randomized QMC, disclosed as such. Survival floors sit at 10^{-300} and the quadrature output is normalized by its sum. The returned map is a reproducible, piecewise-smooth numerical approximation of (2); statements about derivatives refer to derivatives of this approximation, whose fidelity to exact MNP derivatives is governed by L and Q .

Inversion. Given target shares $\hat{p}$ and (V, D) , the inverse transform finds mean-zero μ with $p(\mu) = \hat{p}$ by a damped Jacobi-style diagonal-Newton iteration: all coordinates updated simultaneously using own-location slopes of the unnormalized lattice map (available from the same pass), the field refrozen each iteration, a warm start from the independent-field inverse using total variances, and a tail-aware tolerance (alternatives with target shares below $10^{-4}/N$ are matched best-effort; their locations are weakly identified — itself a conditioning caution for the reduced Jacobian when shares are tiny or alternatives nearly collinear). Convergence in a handful of iterations is an empirical observation on the tested problems, not a theorem. The design follows the original transform’s inversion [10] and the calibration practice of [18].

2 Benchmark against GHK

Anchors first. On the binary case, where MNP has a closed form, the lattice transform agrees to 3.3×10^{-16} and our GHK implementation — validated before being compared against — to 2×10^{-14} at $R = 2 \times 10^5$ (an easy anchor for both methods, useful for code validation rather than as evidence about the hard regime). At $N = 5$ both methods sit within the noise of a 10^7 -draw Monte Carlo reference, and the shipped `thurstone.FactorRace` agrees with the benchmark kernel to 1.1×10^{-5} .

Error metric: throughout, $\max_i |\hat{p}_i - p_i|$ against a simulation reference. A maximum over N coordinates tends to grow with N at fixed per-coordinate accuracy, and the reference’s own maximum-coordinate noise behaves likewise; mean-absolute and total-variation summaries are committed alongside (experiment 16): at $N = 1000$ the lattice’s mean absolute error is 5.6×10^{-6} and total variation 2.8×10^{-3} against a fresh 2×10^6 -draw reference.

N	lattice transform	GHK ($R = 1000$)
5	20 ms, err 5.7×10^{-4}	1 ms, err 3.1×10^{-3}
20	87 ms, err 4.4×10^{-4}	12 ms, err 4.4×10^{-3}
50	214 ms, err 2.7×10^{-4}	75 ms, err 7.9×10^{-3}
200	827 ms, err 3.0×10^{-4}	1275 ms, err 6.8×10^{-3}
1000	4.1 s, err $6.8 \times 10^{-4} \dagger$	—
5000	20.7 s, err $5.0 \times 10^{-4} \dagger$	—

Table 1: Full share vector, $k = 2$ factors, one problem per size (utility spread 1.0 through $N = 200$, 1.5 above). Reference: 2×10^6 MC draws through $N = 200$; 5×10^5 draws at $N = 1000$ and 5000 ($\dagger$: there the reference’s own maximum-coordinate noise is of comparable order, so these entries bound the lattice error only to that order). A replication study at common spread 1.0 (experiment 16; ten problems per size at $N \leq 200$, three at $N = 1000$, twin independent 2×10^6 -draw references per problem) gives worst-case lattice error $3.5\text{--}4.6 \times 10^{-4}$ at every size, at or below the references’ own replicate-to-replicate noise. Lattice error stayed below 10^{-3} throughout; GHK cost crossed over by $N \approx 200$.

Matched-accuracy extrapolation (not a measurement). Fitting a power law to the reference implementation’s timings over $N \in [20, 200]$ ($\alpha \approx 2.0$; a lower bound, the asymptotic per-share cost being cubic) and combining with $R^{-1/2}$ error scaling suggests that matching the lattice’s 5×10^{-4} at $N = 5000$ would take this implementation on the order of forty hours (0.25 h at $R = 1000$, times $(6.8 \times 10^{-3}/5 \times 10^{-4})^2 \approx 185$), against 21 seconds. The baseline is a straightforward single-threaded Python GHK: per-alternative Cholesky, plain pseudorandom uniforms, one realization per problem, shares normalized by their sum. It establishes the scaling of the per-alternative simulation approach, not a field-wide impossibility. The strongest cheap baseline is *direct utility simulation* — draw utilities, take argmax, average — which produces the whole share vector per draw and, unlike GHK, does not pay per alternative. We ran it at wall time matched to the lattice call (experiment 16): its maximum-coordinate error was $1.6\text{--}6\times$ the lattice’s at every tested N (6.1×10^{-4} versus 3.9×10^{-4} at $N = 1000$, 0.5M draws in the matched 4.3 seconds). At the 10^{-3} accuracy level, then, raw forward speed does not by itself separate the methods. What does is error decay, and this is now measured rather than asserted (experiment 17, $N = 200$, $k = 2$): holding the factor rule fixed, lattice self-convergence is 2.7×10^{-8} at $L = 101$ and at machine precision from $L \approx 200$ — the integrand is analytic and endpoint-decaying, so the uniform lattice is effectively spectrally accurate and L -error is subdominant; holding L fixed, the Gauss–Hermite sweep decays from 5.3×10^{-4} (order 3) to 5.9×10^{-9} (order 15); the pre-normalization defect $|1 - \mathbf{1}^\top \tilde{p}|$ is 2.4×10^{-8} . On the same problem the measured accuracy–time frontier (Figure 1, left) has the lattice at 4.1×10^{-5} error in 0.03 seconds (at the twin-reference noise floor), direct MC at 1.2×10^{-4} in 15.7 seconds ($R = 10^7$), scrambled-Sobol direct QMC at 2.0×10^{-4} in 2.6 seconds (2^{20} points), and GHK at 2.0×10^{-3}

in 10.1 seconds ($R = 10^4$): at accuracy targets below 10^{-3} the deterministic map separates decisively, consistent with $R^{-1/2}$ against a rapidly-converging quadrature. Remaining stronger baselines — QMC-GHK, minimax tilting [9] — are queued in the repository.




Figure 1: Left: measured accuracy–time frontier at $N = 200$, $k = 2$ (experiment 17) — every method timed on the same problem against twin 5×10^7 -draw references; direct simulation is a first-class baseline. Right: fixed-design smoothness along a utility path. Fresh-draw GHK shows run-to-run simulation noise (roughness statistic — maximum second difference over dt^2 — 39.3, versus 0.01 for both common-random-number GHK and the lattice); CRN-GHK is smooth conditional on its draw set, and the offset between its curve and the lattice’s is a discrepancy between two approximations whose decomposition requires an independent high-accuracy reference, queued. All three are approximations.

Derivatives, stated correctly. For fixed factor nodes and a fixed lattice, the transform is a reproducible, (piecewise-)differentiable function of μ whose derivatives are evaluated without resampling; analytic own-location slopes of the unnormalized map come from the same pass and serve as the inversion preconditioner. These are derivatives of the numerical approximation, not exact MNP derivatives, and GHK-based practice likewise differentiates its approximation, analytically or with fixed draws [3, 4]. The practical distinction is the character and controllability of the approximation error — quadrature resolution versus draw count — not the possibility of differentiation.

Share inversion. At $N = 1000$, $k = 2$, with targets from 5×10^6 MC draws floored at 10^{-7} and renormalized: the inversion completes in 57 seconds. The iteration solves its own equations to 6.4×10^{-9} — a statement about the solver, not statistical accuracy, since the target carries sampling noise $\sim 2 \times 10^{-4}$ — and recovers the utilities of identified alternatives (target share above 3×10^{-4}) to maximum error 0.017. A replication study (experiment 16) repeats the inversion on a second $N = 1000$ problem against three independent 5×10^6 -draw targets: 57 seconds each, 161–162 alternatives identified (98.1% of share mass), recovery error maximum 0.015–0.023 and median ≈ 0.0027 across replicates, with the recovery-versus-share curve committed alongside. Numerical well-posedness matches Theorem 1: the reduced Jacobian assembled from N JVP calls at $N = 50$ is symmetric to 1.7×10^{-18} , has row-sum defect 6.2×10^{-17} , and reduced eigenvalues in $[2.5 \times 10^{-5}, 0.22]$ — positive definite as proved. We are

not aware of probit share inversion at this scale in current simulation-based practice.

Removal ensemble. All 200 single-removal share vectors from one conditional field construction in $O(QN^2L)$: 51.5 seconds versus 172.4 seconds recomputing ($3.3\times$), with ten evenly spaced removals checked directly (maximum discrepancy 2.8×10^{-17}).

3 Boundary experiments

General covariance through the factor approximation. The method requires $\Sigma \approx VV^\top + D$; GHK does not. Choices depend on (μ, Σ) only through $P\mu$ and $P\Sigma P$, $P = I - \mathbf{1}\mathbf{1}^\top/N$: a common factor loading $V \rightarrow V + \mathbf{1}c^\top$ shifts every conditional location equally and cannot move the argmax (machine-precision fact, tested), so the factor fit must be run in this choice-relevant quotient space — fitting raw Σ can spend scarce rank on an irrelevant common component ($\Sigma = \tau^2\mathbf{1}\mathbf{1}^\top + \beta\beta^\top + D$ with large τ : the raw rank-1 fit takes $\mathbf{1}\mathbf{1}^\top$ and misses $\beta\beta^\top$ entirely; on such an example the contrast-space fit cuts rank-1 share error from 4.6×10^{-2} to 8.3×10^{-3}). A caution from a wrong first implementation: only the common *factor* direction is irrelevant; a common addition to D inflates every difference variance and is choice-relevant, so D must come from the quotient fit. Loadings are centered and SVD-canonicalized (ordered columns, fixed signs), making results reproducible at the covariance level. With this fit: one random correlation matrix at each spectral decay $\lambda_m \propto m^{-\gamma}$, $\gamma \in \{0.5, 1.5, 3\}$, at $N = 50$ (single realizations), share error against an 8×10^6 -draw reference, GHK at the exact Σ for comparison. The plotted quantity is the rank- k approximation error of the stated procedure, not a minimum over factor models. Findings: error decays fastest for near-low-rank truth ($\gamma = 3$: 4.3×10^{-2} at $k = 1$ to 1.5×10^{-3} at $k = 8$) and slowest at intermediate decay ($\gamma = 1.5$: 2.0×10^{-2} to 2.3×10^{-3}). In these three single-matrix illustrations the rank-8 approximation at the stated Q, L had lower observed maximum discrepancy than GHK at $R = 10^4$ (3.2 – 4.1×10^{-3}); this was *not* a wall-time-matched comparison — at this size the $k = 8$ RQMC evaluation takes ~ 14 s against 0.5–0.7 s for GHK at $R = 10^4$ (timings now in the committed output), so at $N = 50$ GHK is the cheaper route to 3 – 4×10^{-3} ; the lattice’s advantage is large N at small k , not this corner. Where accuracy demands exceed the achievable rank- k error, GHK’s arbitrary- Σ generality is the right tool.

Substitution fidelity: a 2×2 factorial. Why want probit-style flexibility at scale? An earlier version of this experiment could not attribute its increments: its Gumbel candidate received factor loadings *and* alternative-specific scales at once (making it a heteroskedastic factor-Gumbel race rather than mixed logit). The redesign isolates the axes. Ground truth is misspecified for every candidate — $t(5)$ factors and skew-normal idiosyncratic noise standardized to mean zero and unit variance, homoskedastic; the race is computed min-wins, so positive skew in race coordinates is *left*-skew in utility coordinates. The four candidates hold the idiosyncratic scale fixed at one and vary $\{\text{Gumbel, Gaussian}\}$ base $\times \{V = 0, V = V^*\}$: independent Luce (= plain multinomial logit exactly), independent probit, factor mixed logit (common-scale, the term now used correctly), and factor probit. All are calibrated to the same full-menu shares (residuals $\leq 4 \times 10^{-10}$). Deletion truths come from 2×10^7 *common* utility draws with the top three alternatives retained per draw, from which every single-deletion and pair-deletion truth follows with common random numbers: 105 informative blocks stratified by deleted mass (26 /

29 / 28 blocks at $>10\%$ / $2\text{--}10\%$ / $0.5\text{--}2\%$), against 24 random blocks before. Mean misallocated fraction of redistributed mass:

	mass $> 10\%$	$2\text{--}10\%$	$0.5\text{--}2\%$
independent Luce (= plain logit)	0.229	0.245	0.261
independent probit	0.210	0.232	0.230
factor mixed logit	0.114	0.146	0.199
factor probit	0.045	0.062	0.094

The factorial decomposition (top two strata): the factor increment is dominant in both families ($+0.115/+0.099$ within Gumbel, $+0.165/+0.170$ within Gaussian); the family increment alone is small ($+0.019/+0.013$ at $V = 0$) but grows with factors present ($+0.069/+0.084$), a negative interaction ($-0.050/-0.071$): factor structure and the Gaussian base are complements in this design. Caveats: a single truth design and seed (replication across skewness, tails, and a separate heteroskedasticity axis is queued), and the skew-normal truth is Gaussian-adjacent. The figure shows individual block observations, not bars.




Figure 2: Left: rank- k contrast-space approximation error under spectral decay γ (dotted: GHK at $R = 10^4$, per γ ; not wall-time matched). Right: the 2×2 factorial substitution study — every informative single-deletion block shown individually.

4 Related work

GHK and simulated likelihood are standard [2, 3, 4]; analytic-approximation lines include MACML [7]; modern Gaussian-probability methods include minimax tilting [9]. Factor dimension reduction for probit is classical: Butler–Moffitt one-factor quadrature [5], Elrod–Keane factor-analytic probit [6], and the probit probability approximations surveyed in [8]. The lattice survival-product identity and its inverse transform are from [10]. On inversion, the social-surplus identity $p = \nabla G$ and convex-conjugate formulation of the inverse are established discrete-choice results — Chiong, Galichon and Shum [15] formulate the inverse through the convex conjugate, and Norets and Takahashi [16] establish surjectivity — and factor-structured MNP for large choice sets has been pursued via Bayesian estimation by Loaiza-Maya and Nibbering [17]. The novelty claimed here is accordingly *not* well-posedness of inversion, nor dimensional reduction, both old: it is the scalable deterministic realization of the forward and inverse maps —

all-alternatives assembly in $O(QNL)$, the $O(QNL)$ Jacobian–vector product, and shared-field deletion counterfactuals in $O(QN^2L)$. Share inversion and its acceleration are central to demand estimation [11, 12]; automatic differentiation of logit-based BLP is recent [13]; a 2026 preprint approximates correlated-choice probabilities with equivariant neural networks [14], evidence of current demand for this computation. We are unaware of an automatically differentiable, large-choice-set factor-probit demand implementation built on deterministic share inversion; that precise object is what Section 5 proposes.

5 Discussion

Estimation is the natural next step: with the reduced coordinates of Proposition 3, gradients of an outer likelihood or GMM objective pass through the share-inversion fixed point by implicit differentiation, a candidate route to differentiable probit demand estimation at large N . Estimating (V, D) rather than supplying them raises the usual identification cautions: utility scale normalization, factor rotation, choice-invariant common loading components, and the fact that only difference covariances are identified. Beyond the Gaussian family, the same machinery computes races with arbitrary alternative-specific idiosyncratic densities; a *common-scale* Gumbel base with zero loadings reproduces the Luce model exactly (alternative-specific scales break Luce, as they should), yielding correlated-softmax generalizations. Finally, the inversion inherits the original transform’s interpolation structure (per-node location-to-probability curves are cross-correlations, evaluable at all offsets simultaneously), which we have not yet exploited.

References

- [1] L. L. Thurstone. A law of comparative judgment. *Psychological Review* 34 (1927) 273–286.
- [2] J. Geweke, M. Keane, D. Runkle. Alternative computational approaches to inference in the multinomial probit model. *Review of Economics and Statistics* 76 (1994) 609–632.
- [3] V. Hajivassiliou, D. McFadden. The method of simulated scores for the estimation of LDV models. *Econometrica* 66 (1998) 863–896.
- [4] K. Train. *Discrete Choice Methods with Simulation*, 2nd ed. Cambridge, 2009.
- [5] J. S. Butler, R. Moffitt. A computationally efficient quadrature procedure for the one-factor multinomial probit model. *Econometrica* 50 (1982) 761–764.
- [6] T. Elrod, M. P. Keane. A factor-analytic probit model for representing the market structure in panel data. *Journal of Marketing Research* 32 (1995) 1–16.
- [7] C. Bhat. The maximum approximate composite marginal likelihood (MACML) estimation of multinomial probit-based unordered response choice models. *Transportation Research B* 45 (2011) 923–939.
- [8] R. Connors, S. Hess, A. Daly. Analytic approximations for computing probit choice probabilities. *Transportmetrica A* 10 (2014) 119–139.
- [9] Z. Botev. The normal law under linear restrictions: simulation and estimation via minimax tilting. *JRSS-B* 79 (2017) 125–148.

- [10] P. Cotton. Inferring relative ability from winning probability in multientrant contests. *SIAM J. Financial Mathematics* 12 (2021) 295–317.
- [11] S. Berry, J. Levinsohn, A. Pakes. Automobile prices in market equilibrium. *Econometrica* 63 (1995) 841–890.
- [12] C. Conlon, J. Gortmaker. Best practices for differentiated products demand estimation with pyblp. *RAND Journal of Economics* 51 (2020) 1108–1161.
- [13] A. Chia. Automatically differentiable random coefficient logistic demand estimation. arXiv:2106.04636 (2021).
- [14] E. Huch, M. Keane. Amortized inference for correlated discrete choice models via equivariant neural networks. arXiv:2603.24705 (2026).
- [15] K. Chiong, A. Galichon, M. Shum. Duality in dynamic discrete-choice models. *Quantitative Economics* 7 (2016) 83–115.
- [16] A. Norets, S. Takahashi. On the surjectivity of the mapping between utilities and choice probabilities. *Quantitative Economics* 4 (2013) 149–155.
- [17] R. Loaiza-Maya, D. Nibbering. Scalable Bayesian estimation in the multinomial probit model. *Journal of Business & Economic Statistics* 40 (2022) 1678–1690.
- [18] The `allocation` package: Thurstone-tilt portfolio construction. <https://github.com/microprediction/allocation>.