

Softmax from Fast Mixing

Peter Cotton*

August 2026 — working draft

Abstract

Softmax is used as a choice law almost everywhere, usually without a mechanism that produces it. We give one. When N channels compete under a shared hidden state that mixes faster than the channels fire, the winning probability is softmax in the stationary mean intensities at leading order, and the first correction is a Green–Kubo term built from integrated rate autocovariances. One pair of objects, the mean intensities and the correlation matrix K , then predicts every blocked-subset counterfactual by restricting sums, so the whole deletion ensemble follows from one calculation. Fluctuations common to all channels cancel exactly, so departures from softmax measure differential structure alone, and environmental modes of rank r give K of rank r , up to a cap set by the environment’s own dimension. The same correction predicts the runner-up kernel, which is observable in any race that reports a finishing order, so the theory can be tested without running interventions. We verify the expansion on finite chains, then on Brownian narrow escape with reaction windows, where the leading and corrected errors decay with measured orders 0.93 and 1.93 against an exactly discretized generator.

1 Introduction

A choice among competing options is often modelled by softmax. The rule assigns option i a probability proportional to e^{θ_i} , and it appears in discrete choice (Luce, 1959; McFadden, 1974), reinforcement learning, statistical physics, and the output layer of essentially every classifier (Bridle, 1990). Its usual justification is convenience. Softmax is smooth, normalized, and differentiable, and in the Gumbel random-utility model it is exact, a representation Yellott (1977) showed is the only one consistent with Luce’s axiom under independent discriminial processes.

Underneath many applications, though, there is real kinetics. A chemical reaction fires when one of several channels crosses a barrier first. A diffusing particle escapes through whichever window it encounters first. A kinetic Monte Carlo (KMC) step selects among escape channels by their rates (Bortz et al., 1975; Gillespie, 1977). In these systems the choice is a race, and the race has a mechanism. That a race among competing channels is at once a competing-risks model and a random-utility choice model is an old observation (Marley and Colonius, 1992). It is worth asking what that mechanism implies about the choice law, rather than assuming softmax and moving on.

This paper answers that for one specific and common situation. The channels share a hidden state, and that state mixes fast relative to the rate at which channels fire. Catalysis supplies the

*Code and machine verification for every claim in this paper: <https://github.com/microprediction/kinetics> (experiments 7 and 9). Web companion: <https://kinetics.microprediction.org>.

picture. Adsorbate configurations rearrange quickly while the reaction event is rare, so every channel sees an environment that has effectively averaged out. Chemical physics calls this regime dynamical disorder (Zwanzig, 1990), and the same logic underlies stochastically gated reactions (Szabo et al., 1982). Under that separation we show two things.

The leading-order result is that softmax is exact in the limit, with $\theta_i = \log \bar{\lambda}_i$ the log of the stationary mean intensity. Fast mixing homogenizes the environment, each channel is left with its average rate, and the race among constant rates is the exponential race, whose choice law is Luce (Cox, 1959).

The first correction is a Green–Kubo term. It is built from $K_{jk} = \int_0^\infty \text{Cov}_\pi(\lambda_j(Y_0), \lambda_k(Y_t)) dt$, the time-integrated covariance of the firing intensities under the stationary law. This is a dynamical object rather than a static one. Two channels whose rates are correlated at equal times contribute nothing unless the correlation persists over the timescale on which the environment decorrelates.

The correction is what breaks independence of irrelevant alternatives (IIA). Softmax redistributes a removed channel’s probability in proportion to the survivors’ probabilities, which is the IIA prediction. Where it has been measured, that prediction fails. Diversion ratios in empirical demand work depart systematically from proportionality (Conlon and Mortimer, 2021), shared features break IIA in choice experiments (Tversky, 1972), and the standard specification test for logit is built on exactly this restriction (Hausman and McFadden, 1984). The Green–Kubo term supplies the deviation, and it does so for every blocked subset at once.

That last point is the practical one. The pair $(\bar{\lambda}, K)$ is computed once. Any counterfactual, remove channel 3, block channels $\{2, 5, 7\}$, is answered by restricting the sums in the correction to the surviving set. No new solve is required per counterfactual. This is the statistical counterpart of a theme running through our related work, that one global computation can contain an entire leave-one-out ensemble.

Two structural facts sharpen the result. Fluctuations shared by all channels, $\lambda_i(y) = a_i c(y)$, cancel from the correction identically. We show this is not merely an asymptotic statement. A common fluctuation is a common time change, so Luce holds exactly at every degree of separation, which is the proportional-hazards normal form. Separately, if the intensity deviations are driven by r environmental modes then K has rank r , so the correction inherits any low-dimensional structure the physics has.

Section 2 sets up the shared-state race and the killed resolvent. Section 3 states and proves the expansion. Section 4 collects the three structural properties. Section 5 verifies everything on finite chains, where each quantity is an exact linear solve. Section 6 moves to continuous physics, a reflected Brownian motion in a disk with reaction windows, and reports the measured convergence orders together with two systematic errors found there. Section 7 shows the correction also predicts the runner-up kernel, which makes it testable against ranked data.

2 The shared-state race

Let Y_t be an ergodic Markov process on a state space with generator L/ε and invariant law π . The parameter ε is the ratio of the environment’s mixing time to the timescale on which channels fire. Small ε means fast mixing.

Each channel $i \in \{1, \dots, N\}$ fires at intensity $\lambda_i(Y_t) \geq 0$, conditionally Poisson given the path of Y . The first channel to fire wins. The model is a marked Markov-modulated Poisson process (Fischer and Meier-Hellstern, 1993) observed until its first event, and the λ_i are cause-

specific hazards in the competing-risks sense. For an available set $A \subseteq \{1, \dots, N\}$ write $\Lambda_A(y) = \sum_{j \in A} \lambda_j(y)$ for the total hazard, and write $\bar{\lambda}_i = \int \lambda_i d\pi$ for the stationary mean intensity.

Let $u_i^A(y)$ be the probability that channel i wins, started from $Y_0 = y$. Conditioning on the first instant and using the conditional Poisson structure gives

$$\left(\frac{L}{\varepsilon} - \Lambda_A\right) u_i^A = -\lambda_i, \quad u_i^A = \left(\Lambda_A - \frac{L}{\varepsilon}\right)^{-1} \lambda_i. \quad (1)$$

This is a killed resolvent. The generator propagates the environment and the multiplication operator Λ_A kills paths at the total firing rate. On a finite state space it is a linear system, which is what makes the verification in Section 5 exact.

The observable of interest is the win probability from a stationary start,

$$p_i^A = \int u_i^A d\pi. \quad (2)$$

Summing (1) over $i \in A$ gives $\sum_{i \in A} u_i^A = 1$, so the p_i^A are a probability vector on A for every ε .

3 Homogenization and the Green–Kubo correction

The expansion below is the averaging principle for a fast ergodic driver (Khas'minskii, 1966) in the operator form of Kurtz (1973), and its solvability conditions are the usual cell problem (Pavliotis and Stuart, 2008, Ch. 9). Write $\tilde{\lambda}_i = \lambda_i - \bar{\lambda}_i$ for the fluctuation of the i th intensity, and Π for the projection $\Pi f = (\int f d\pi) \mathbf{1}$ onto constants. The deviation operator

$$D = (\Pi - L)^{-1}(I - \Pi) \quad (3)$$

inverts L on the mean-zero subspace, and $Dg = \int_0^\infty e^{tL} g dt$ for mean-zero g . Note the sign it carries. Differentiating that integral gives $LDg = -g$, so D is minus the pseudo-inverse of L , and the sign propagates into the proof below. The Green–Kubo matrix is

$$K_{jk} = \int_0^\infty \text{Cov}_\pi(\lambda_j(Y_0), \lambda_k(Y_t)) dt = \int \tilde{\lambda}_j(D\tilde{\lambda}_k) d\pi. \quad (4)$$

The second expression is what one computes. It replaces a time integral by a single linear solve. Note that K need not be symmetric. It is symmetric when the environment is reversible, and the verification in Section 5 deliberately uses a non-reversible chain so that the index placement in the theorem is actually tested.

The diagonal of K is not new. For a single channel with a fluctuating rate the integral $\int_0^\infty \text{Cov}_\pi(\lambda(Y_0), \lambda(Y_t)) dt$ is the classical second-order cumulant correction to an effective decay rate, known since Anderson (1954) and Kubo (1954) as motional narrowing and developed as a general expansion by van Kampen (1974). The name Green–Kubo is used here in the sense of Green (1954) and Kubo (1957), an integrated dynamical correlation standing in for a transport coefficient.

What that scalar theory cannot see is the object this paper is about. A rate correction common to all channels cancels from a ratio, so it leaves the choice law untouched. The branching ratio is moved instead by the off-diagonal and generally non-symmetric entries K_{ji} , and those are what Theorem 1 puts to work.

Theorem 1 (Homogenized choice law and its first correction). *Let $\bar{\Lambda}_A = \sum_{j \in A} \bar{\lambda}_j$ and suppose $\bar{\Lambda}_A > 0$. Then as $\varepsilon \rightarrow 0$,*

$$p_i^A = \frac{\bar{\lambda}_i}{\bar{\Lambda}_A} - \frac{\varepsilon}{\bar{\Lambda}_A} \left[\sum_{j \in A} K_{ji} - \frac{\bar{\lambda}_i}{\bar{\Lambda}_A} \sum_{j,k \in A} K_{jk} \right] + O(\varepsilon^2). \quad (5)$$

In particular $p_i^A \rightarrow \bar{\lambda}_i / \bar{\Lambda}_A$, the softmax law in the log mean intensities $\theta_i = \log \bar{\lambda}_i$.

Proof. Multiply (1) by ε to get $(L - \varepsilon \Lambda_A)u_i^A = -\varepsilon \lambda_i$, expand $u_i^A = u_i^{(0)} + \varepsilon u_i^{(1)} + \varepsilon^2 u_i^{(2)} + O(\varepsilon^3)$, and collect powers of ε .

At order ε^0 , $Lu_i^{(0)} = 0$, so $u_i^{(0)}$ is a constant c_i by ergodicity.

At order ε^1 the equation reads $Lu_i^{(1)} = \Lambda_A c_i - \lambda_i$. A solution exists only if the right side is orthogonal to π , which forces $\bar{\Lambda}_A c_i = \bar{\lambda}_i$ and gives the leading term. The right side is then $c_i \tilde{\Lambda}_A - \tilde{\lambda}_i$ with $\tilde{\Lambda}_A = \sum_{j \in A} \tilde{\lambda}_j$, and inverting with $LD = -I$ on mean-zero functions gives

$$u_i^{(1)} = -D(c_i \tilde{\Lambda}_A - \tilde{\lambda}_i) + m_i,$$

where the constant $m_i = \int u_i^{(1)} d\pi$ is fixed at the next order.

At order ε^2 the equation is $Lu_i^{(2)} = \Lambda_A u_i^{(1)}$, whose solvability condition is $\int \Lambda_A u_i^{(1)} d\pi = 0$. The constant part of Λ_A contributes $\bar{\Lambda}_A m_i$, the constant m_i integrates against $\tilde{\Lambda}_A$ to zero, and what remains is

$$0 = \bar{\Lambda}_A m_i - \int \tilde{\Lambda}_A D(c_i \tilde{\Lambda}_A - \tilde{\lambda}_i) d\pi.$$

Using (4), $\int \tilde{\Lambda}_A D\tilde{\lambda}_i d\pi = \sum_{j \in A} K_{ji}$ and $\int \tilde{\Lambda}_A D\tilde{\Lambda}_A d\pi = \sum_{j,k \in A} K_{jk}$, so

$$m_i = -\frac{1}{\bar{\Lambda}_A} \left[\sum_{j \in A} K_{ji} - \frac{\bar{\lambda}_i}{\bar{\Lambda}_A} \sum_{j,k \in A} K_{jk} \right].$$

Finally $p_i^A = c_i + \varepsilon m_i + O(\varepsilon^2)$ by (2), which is (5). $\square$

Proposition 1 (General start). *Let the environment start from an arbitrary law μ_0 rather than from π , and write $p_i^A(\mu_0) = \int u_i^A d\mu_0$. Then*

$$p_i^A(\mu_0) = \frac{\bar{\lambda}_i}{\bar{\Lambda}_A} + \varepsilon \left[m_i - \int D\left(\frac{\bar{\lambda}_i}{\bar{\Lambda}_A} \tilde{\Lambda}_A - \tilde{\lambda}_i\right) d\mu_0 \right] + O(\varepsilon^2), \quad (6)$$

with m_i the constant appearing in the proof of Theorem 1. The leading softmax term does not depend on the start. The extra term vanishes when $\mu_0 = \pi$, because D maps into mean-zero functions, which recovers (5).

Proof. The expansion in the proof of Theorem 1 is pointwise in y , and integrating $u_i^{(1)} = -D(c_i \tilde{\Lambda}_A - \tilde{\lambda}_i) + m_i$ against μ_0 instead of π leaves the displayed extra term. $\square$

This is worth stating because it is easy to apply (5) to a system observed away from equilibrium and silently lose the correction. Experiment 40 measures the loss. From a point-mass start the stationary formula (5) is accurate only to order 0.968, while (6) restores order 1.963. Under the invariant law the extra term is 7.3×10^{-17} and the two forms agree.

Remark 1. The two terms in the bracket have distinct roles. The first, $\sum_{j \in A} K_{ji}$, measures how channel i 's rate correlates dynamically with the total hazard it competes against. The second subtracts the part of that effect which is common to all channels, and is what keeps $\sum_{i \in A} p_i^A = 1$ at each order.

4 Three structural properties

One computation answers every counterfactual. The pair $(\bar{\lambda}, K)$ does not depend on A . Every blocked-set counterfactual is obtained from (5) by restricting the sums to the surviving set. For a system with N channels this reduces 2^N separate resolvent solves to one deviation solve and $O(N^2)$ stored numbers, at the cost of the $O(\varepsilon^2)$ error the expansion carries.

Common-mode fluctuations cancel. Suppose $\lambda_i(y) = a_i c(y)$, so that every channel is modulated by the same environmental factor. Then $\bar{\lambda}_i = a_i \bar{c}$ and $K_{jk} = a_j a_k \gamma$ with $\gamma = \int \bar{c} (D\bar{c}) d\pi$. Substituting into the bracket in (5) gives $a_i \gamma \sum_{j \in A} a_j - (a_i / \sum_{j \in A} a_j) \gamma (\sum_{j \in A} a_j)^2 = 0$. The correction vanishes identically.

Proposition 2 (Proportional hazards give Luce exactly). *Suppose $\lambda_i(y) = a_i c(y)$ with $a_i > 0$ and $c \geq 0$ not identically zero. Then for every y , every ε , and every A ,*

$$u_i^A(y) = \frac{a_i}{\sum_{j \in A} a_j},$$

so the win probability is Luce's law from any start, not merely from π and not merely to the order computed above.

Proof. Let v_i be the constant function $a_i / \sum_{j \in A} a_j$. Constants are annihilated by L , and $\Lambda_A v_i = c(y) (\sum_{j \in A} a_j) \cdot a_i / \sum_{j \in A} a_j = a_i c(y) = \lambda_i$, so $(\Lambda_A - L/\varepsilon) v_i = \lambda_i$. Since the killed resolvent is invertible, v_i is the unique solution of (1), which is u_i^A . $\square$

A common multiplicative modulation is a shared frailty in the survival sense (Vaupel et al., 1979), and the invariance it produces is the proportional-hazards structure of Cox (1972). The proof needs no positivity of c , no time change, and no condition on ε . What it says is that a common modulation factors out of the competition entirely, which is the proportional-hazards normal form. Experiment 40 confirms it where the hypotheses are stressed. With c vanishing on four of six states, with c ranging over nine orders of magnitude, and at $\varepsilon = 100$, the departure from Luce stays below 5.5×10^{-14} at every state.

This is a stronger statement than the vanishing of the correction, and it is practically useful in the opposite direction. It says a measured departure from softmax cannot be produced by any amount of shared environmental fluctuation. Whatever deviation is observed is evidence of differential structure among the channels. It is worth setting this against the classical negative results. The latent joint law of competing risks is not identified from what one observes (Tsiatis, 1975), and identification must be bought with extra structure such as regressor variation (Heckman and Honoré, 1989). The statement here runs the other way. One large class of environmental dependence, the perfectly shared one, leaves no trace whatever in the choice law, so any trace that does appear rules that class out.

Low-rank environments give low-rank corrections. Suppose the intensity deviations are driven by r environmental modes, $\lambda_i - \bar{\lambda}_i = b_i^\top z(y)$ with z taking values in $\mathbb{R}^r$ and $\int z d\pi = 0$. Then $K = B\Gamma B^\top$ with B the $N \times r$ matrix of loadings and $\Gamma_{ab} = \int z_a(Dz_b)d\pi$, so K has rank at most r . There is a second cap. Deviations live in the mean-zero subspace of the environment, so $\text{rank}(K) \leq \min(N, m-1)$ for an environment with m states, whatever the loadings are. Experiment 40 attains both bounds, giving rank exactly r for $r = 1, \dots, 5$ on a six-state environment and rank $5 = \min(10, 5)$ for generic intensities on ten channels. The compressed correlated race that appears elsewhere in this program therefore has a physical provenance rather than being a modelling convenience. The same low-dimensional error-components structure is imposed by hand in mixed logit (McFadden and Train, 2000), where it is a specification choice. Here it is inherited from the environment.

5 Verification on finite chains

On a finite state space every object above is an exact linear solve, so the theorem can be tested without simulation error. Experiment 7 runs a six-state chain with ten channels.

Setup. The generator has independent unit-mean exponential off-diagonal rates, giving a non-reversible chain with spectral gap 2.63. Channel intensities are lognormal with log-scale 0.8, giving $\bar{\Lambda}_A = 11.4$ on the full set. The seed is 5 and the sweep covers thirteen values of ε from 10^{-3} to 1. The dimensionless small parameter is $\varepsilon \bar{\Lambda}_A / \text{gap} \approx 4.3\varepsilon$, so the fitted window ends near 0.08 in those units and $\varepsilon = 0.3$ corresponds to 1.3.

The measured convergence order of the leading term is 0.985, against the theoretical 1. After adding the Green–Kubo correction the measured order is 1.985, against the theoretical 2. Both are fitted over the smallest six of thirteen values of ε , spanning 10^{-3} to 1.8×10^{-2} , since the expansion is asymptotic and the largest ε carry the $O(\varepsilon^2)$ residual. The gain of exactly one order is the signature the expansion predicts.

One instance is an anecdote, so experiment 40 repeats the measurement over twenty environments drawn with four to nine hidden states, three to eight channels, and intensity dispersions from 0.3 to 1.2. The leading order has median 0.960 and range $[0.919, 0.982]$, the corrected order has median 1.958 and range $[1.888, 1.979]$, and the gain exceeds 0.948 in every single trial.

Figure 1 shows both. The subset-uniformity claim is tested by fixing one $(\bar{\lambda}, K)$ pair and predicting every availability set. Experiment 40 sweeps all 1013 subsets of size at least two at $\varepsilon = 0.05$, so the maxima here are true maxima rather than samples. The worst error over all subsets falls from 3.1×10^{-3} under softmax to 4.4×10^{-4} with the correction, with no per-subset computation.

The order is uniform in a strong sense. Fitting the convergence order separately on each of the 1013 subsets puts every corrected order in $[1.92, 2.00]$. Small subsets are better rather than worse, because the effective parameter is $\varepsilon \bar{\Lambda}_A / \text{gap}$ and $\bar{\Lambda}_A$ shrinks with the subset. Exactly one subset has the correction slightly worse than softmax, a two-element near-tie where both errors are already two orders of magnitude below the maximum.

The word Green–Kubo is doing work in (4), and experiment 40 tests whether it needs to. Replacing K by the equal-time covariance $\int \tilde{\lambda}_j \tilde{\lambda}_k d\pi$ scaled by a single timescale, with that timescale fitted to give the ablation its best case, leaves the measured order at 1.09 rather than 1.96, and the error at the smallest ε is 8.8 times worse. Static correlation between two channels does not determine the correction. The time integral does.




Figure 1: Finite-chain verification (experiment 7). Left, error against ε for the leading softmax law and for the Green–Kubo corrected law, with measured orders 0.985 and 1.985. Right, the same $(\bar{\lambda}, K)$ pair applied to 41 random availability sets at $\varepsilon = 0.05$, with no per-subset computation.

Common-mode cancellation is confirmed at 2.6×10^{-16} for the Green–Kubo coefficient, and the exact win probabilities agree with Luce to 1.1×10^{-16} at $\varepsilon = 0.3$. The regression suite carries the same check at $\varepsilon = 3.0$, where the environment is ten times slower than the channels rather than faster and the expansion has no claim at all, and Luce still holds to 10^{-12} . That is the content of Proposition 2. Rank-two environmental loadings give K with singular values $[0.022, 0.002, 0, 0]$, so the rank is exactly two.

A Gillespie simulation of the same system, event-driven and including the hidden state, agrees with the resolvent solution to 1.9×10^{-3} across 40,000 trajectories at $\varepsilon = 0.3$, against a binomial sampling noise of 2.1×10^{-3} .

The chain is not reversible, so K is asymmetric. On this instance $\max_{jk} |K_{jk} - K_{kj}| = 0.139$ against $\max_{jk} |K_{jk}| = 0.603$, a relative asymmetry of 23% recorded by experiment 40. This is what distinguishes K_{ji} from K_{ij} in (5), and experiment 40 confirms the index order is not a convention. Transposing K in the correction degrades the measured order from 1.94 to 1.58, and K agrees with brute-force time integration of the covariance to 4.1×10^{-6} .

6 Narrow escape

Finite chains verify the algebra. To test the theorem as a statement about physics we take the hidden state to be reflected Brownian motion in the unit disk and the channels to be reaction windows on the boundary, bands in which channel i fires at rate κ . This is the narrow escape geometry (Schuss et al., 2007) with partially reactive patches, and the leading band-area law is the well-mixed limit of Berg and Purcell (1977). The first window to fire wins. Here κ plays the role of ε , since a slow reaction rate means the particle mixes many times before firing.

Setup. Sixteen boundary windows on the unit circle, six clustered within about 50 degrees and the rest scattered, drawn with seed 42. Angular half-widths are lognormal with median 0.030,

κ	softmax	with Green-Kubo	improvement
2	1.6×10^{-3}	5.0×10^{-5}	31×
5	3.7×10^{-3}	3.0×10^{-4}	12×
10	6.9×10^{-3}	1.1×10^{-3}	6.3×
20	1.2×10^{-2}	3.9×10^{-3}	3.2×
40	2.0×10^{-2}	1.2×10^{-2}	1.6×

Table 1: Maximum absolute error in the win probabilities against the exactly discretized generator, over sixteen boundary windows. The correction gains an order of magnitude at small κ and degrades to nothing by $\kappa = 40$, which is where the neglected $O(\kappa^2)$ term takes over.

ranging over $[0.019, 0.055]$ on this draw, and three pairs overlap. The reaction bands extend 0.12 inward in radius. The polar grid is 40×88 , giving 3520 cells. The resolvent sweep covers $\kappa \in \{2, 5, 10, 20, 40\}$, and the Brownian simulation runs at $\kappa \in \{5, 20\}$ with time step 2×10^{-4} and 50,000 trajectories.

Experiment 9 separates three layers so that no error source can hide in another. Layer A is a real Brownian simulation with independent per-step firing. Layer B is an exact killed-resolvent solve on a finite-volume Neumann Laplacian over a polar grid, with the uniform invariant law, reversibility, and conservation checked numerically. Layer C is the theory, band-area softmax plus the Green-Kubo term from one deviation solve, using the same code that passed the chain tests.

Two homogenizations should be kept apart here. Replacing a patchy boundary by an effective uniform reactivity is boundary homogenization (Berezhkovskii et al., 2004; Lawley and Keener, 2015), and expansions in a small geometric parameter give splitting probabilities among small targets (Bressloff, 2020). The small parameter in this paper is temporal rather than geometric, so the two expansions are orthogonal and compose.

Comparing layers B and C tests the theorem on a continuous-physics generator. The softmax error decays with measured order 0.93 and the corrected error with order 1.93, again a clean gain of one order. Both orders are fitted on the three smallest values of κ , for the same reason as before.

Table 1 gives the full sweep, including where the correction stops paying. The improvement factor falls monotonically from 31× at $\kappa = 2$ to 1.6× at $\kappa = 40$. This is the expected behaviour of a first-order expansion rather than a defect, but it bounds the useful range, and a practitioner should read the correction as buying roughly an order of magnitude only while κ stays in the first half of this table.

Figure 2 shows the three layers together. Comparing layers A and B tests the discretization. The maximum discrepancy is 2.2×10^{-3} at $\kappa = 5$ and 1.6×10^{-3} at $\kappa = 20$, against a sampling noise of about 10^{-3} at 50,000 trajectories.

Two systematic errors. Both would have corrupted a direct comparison of simulation against theory, and neither is visible without the middle layer.

The first is band quantization, worth 5×10^{-2} . Sharp cell indicators mis-size any window narrower than a grid cell, with the mis-sizing ratio running from 2.0 down to zero, since the narrowest window here spans about half a cell and is missed entirely. Fractional area-exact cell coverage fixes it, after which band areas match their analytic values to 2×10^{-17} .




Figure 2: The theorem on continuous physics (experiment 9). Reflected Brownian motion in the unit disk with boundary reaction windows, error against κ for the leading and corrected laws, measured orders 0.93 and 1.93 against an exactly discretized Neumann generator.

The second is overlap semantics, worth 1.4×10^{-2} . The window geometry contains three overlapping pairs. The original simulation awarded an overlap to the lower-indexed window, while the model stacks intensities as $\lambda = \kappa \sum_i g_i$, which is the physically correct convention for independent receptors. The error's structure identified it, being -0.014 on one window and $+0.006$ on its overlap partner, and its flatness under both time-step and grid refinement confirmed that it was not a discretization artifact.

7 The correction predicts the runner-up

The correction has been described so far as a statement about win probabilities, which are what a blocked-set experiment measures. Those experiments are expensive, since each availability set requires its own intervention. There is a cheaper observable that carries the same information.

Let $M_{ij} = \mathbb{P}(\text{runner-up} = j \mid \text{winner} = i)$ be the runner-up kernel, the probability that j finishes second given that i finishes first. Removing channel i changes the outcome only in those realizations where i won, and in those the runner-up inherits the win. This gives an identity that holds for every ε and involves no approximation,

$$q_j^{(-i)} = p_j + p_i M_{ij}, \quad (7)$$

where $q^{(-i)}$ is the deletion counterfactual. The companion work shows M is exactly the object that identifies single deletions, and that winner-only data constrains it not at all.

Applying Theorem 1 to the full set and to $A \setminus \{i\}$ therefore predicts M itself. At leading order the prediction is

$$M_{ij} = \frac{\bar{\lambda}_j}{\bar{\Lambda}_A - \bar{\lambda}_i} + O(\varepsilon), \quad (8)$$

which says the runner-up is drawn from the race among the remaining channels, the IIA prediction. Substituting the corrected shares into (7) gains an order.

Experiment 39 verifies this on the same chains. The identity (7) holds to 1.8×10^{-15} , the exactly computed kernel is row-stochastic to 8.0×10^{-15} , which checks the two-stage solve, and the measured convergence orders are 0.967 for the Luce prediction (8) and 1.967 with the Green-Kubo correction. At $\varepsilon = 0.005$ the maximum error over off-diagonal entries falls from 3.0×10^{-4} to 2.1×10^{-6} .

This matters for testing the theory. The runner-up is observed in any race that reports a finishing order, so the correction can be checked against ordinary ranked data rather than against interventions. It also runs the other way. Where the environment is unknown, a measured M constrains the Green–Kubo matrix, which is the beginning of estimating K from data rather than computing it from a known generator.

8 Reproducibility

Every number in this paper comes from a committed script. Experiment 7 gives the chain verification, experiment 9 the narrow-escape study, experiment 39 the runner-up bridge, and experiment 40 the replication, the equal-time ablation, the rank bounds, and the index-order check. All four run from a clean checkout of <https://github.com/microprediction/kinetics> with NumPy, SciPy, and Matplotlib, and each writes the `results.csv` the corresponding numbers are quoted from. Experiments 7, 39, and 40 take seconds. Experiment 9 takes about three minutes. A regression suite locks every property stated here, including the two convergence orders, the exactness of common-mode cancellation, the rank bounds, and the runner-up identity.

9 Related work

The averaging that produces the leading term is classical (Khas’minskii, 1966; Kurtz, 1973; Pavliotis and Stuart, 2008; Yin and Zhang, 1998), and the closest recent instance is Colantoni (2026), who studies a killing rate modulated by a continuous-time Markov chain, gives Feynman–Kac representations, and proves an averaging principle in the fast-switching limit. That work stops where this one starts. It carries a single killing rate rather than competing channels, so there is no branching ratio to correct, and it establishes the limit without the first-order term.

The ingredient is classical on the other side too. Motional narrowing (Anderson, 1954; Kubo, 1954; van Kampen, 1974) and dynamical disorder (Zwanzig, 1990; Szabo et al., 1982) have used integrated rate autocovariances for seventy years, but for effective rates and survival probabilities, in the scalar setting where no choice law arises. Recent linear-response work on first-passage problems predicts the response of the mean first-passage *time* rather than the identity of the winner.

For marked Markov-modulated processes there are matrix-analytic methods (Fischer and Meier-Hellstern, 1993) that compute mark probabilities exactly, with no expansion and no asymptotic regime. Theorem 1 is the fast-environment asymptotic of those same quantities, and what it adds is not the numbers but the structure the asymptotic exposes. One pair $(\bar{\lambda}, K)$ serves every subset, shared modulation drops out, and low-rank environments give low-rank corrections.

That a race is at once a competing-risks model and a random-utility model is old (Marley and Colonius, 1992), as is the race-to-Luce mechanism in other fields. What appears not to have been stated is that softmax is the homogenized choice law of fast hidden kinetics, with a computable first correction.

10 Discussion

The result explains why softmax is a reasonable default in systems with fast shared dynamics, and it quantifies exactly when that default fails. The failure is governed by integrated dynamical

correlations, not by static ones, and it is insensitive to any fluctuation shared across channels.

Four limitations bound the claim. The expansion is in the mixing-to-firing timescale ratio and says nothing when channels fire faster than the environment relaxes, which is the quenched limit of dynamical disorder (Zwanzig, 1990). The correction is first order, so systems at moderate separation retain an $O(\varepsilon^2)$ residual, quantified in Table 1, where the improvement decays to nothing by $\kappa = 40$. The theorem is proved for finite chains, and a formal treatment of reversible diffusions with a spectral gap is left open, so the narrow-escape study verifies rather than proves that case. The correction as usually quoted assumes the environment is observed in equilibrium, and Proposition 1 gives the extra term otherwise. Finally the continuum study is one geometry with one window arrangement, and while the chain results replicate over twenty environments, the physics results do not yet replicate over geometries.

One extension looks tractable. The disk’s boundary operator diagonalizes in Fourier modes, which may render the leading geometric correction analytic rather than numerical.

Section 7 closes the gap between the computed correction and the estimated one, at least in one direction. What remains is the inverse problem. Estimating K from an observed runner-up kernel is an underdetermined fit, since M supplies $N(N - 1)$ numbers for the N^2 entries of K and only through the combinations appearing in (5), so the identifiable part of K has yet to be characterized.

References

- P. W. Anderson. A mathematical model for the narrowing of spectral lines by exchange or motion. *Journal of the Physical Society of Japan* 9 (1954) 316–339.
- A. M. Berezhkovskii, Yu. A. Makhnovskii, M. I. Monine, V. Yu. Zitserman, S. Y. Shvartsman. Boundary homogenization for trapping by patchy surfaces. *Journal of Chemical Physics* 121 (2004) 11390–11394.
- H. C. Berg, E. M. Purcell. Physics of chemoreception. *Biophysical Journal* 20 (1977) 193–219.
- A. B. Bortz, M. H. Kalos, J. L. Lebowitz. A new algorithm for Monte Carlo simulation of Ising spin systems. *Journal of Computational Physics* 17 (1975) 10–18.
- P. C. Bressloff. Asymptotic analysis of extended two-dimensional narrow capture problems. arXiv:2010.11885 (2020).
- J. S. Bridle. Probabilistic interpretation of feedforward classification network outputs, with relationships to statistical pattern recognition. In *Neurocomputing*, NATO ASI Series F 68 (1990) 227–236.
- F. Colantoni. Time-dependent Robin heat equation via Markovian switching. arXiv:2604.27901 (2026).
- C. Conlon, J. H. Mortimer. Empirical properties of diversion ratios. *RAND Journal of Economics* 52 (2021) 693–726.
- D. R. Cox. The analysis of exponentially distributed life-times with two types of failure. *Journal of the Royal Statistical Society B* 21 (1959) 411–421.

- D. R. Cox. Regression models and life-tables. *Journal of the Royal Statistical Society B* 34 (1972) 187–202.
- W. Fischer, K. Meier-Hellstern. The Markov-modulated Poisson process (MMPP) cookbook. *Performance Evaluation* 18 (1993) 149–171.
- D. T. Gillespie. Exact stochastic simulation of coupled chemical reactions. *Journal of Physical Chemistry* 81 (1977) 2340–2361.
- M. S. Green. Markoff random processes and the statistical mechanics of time-dependent phenomena II. Irreversible processes in fluids. *Journal of Chemical Physics* 22 (1954) 398–413.
- J. Hausman, D. McFadden. Specification tests for the multinomial logit model. *Econometrica* 52 (1984) 1219–1240.
- J. J. Heckman, B. E. Honoré. The identifiability of the competing risks model. *Biometrika* 76 (1989) 325–330.
- R. Z. Khas'minskii. A limit theorem for the solutions of differential equations with random right-hand sides. *Theory of Probability and Its Applications* 11 (1966) 390–406.
- R. Kubo. Note on the stochastic theory of resonance absorption. *Journal of the Physical Society of Japan* 9 (1954) 935–944.
- R. Kubo. Statistical-mechanical theory of irreversible processes I. *Journal of the Physical Society of Japan* 12 (1957) 570–586.
- T. G. Kurtz. A limit theorem for perturbed operator semigroups with applications to random evolutions. *Journal of Functional Analysis* 12 (1973) 55–67.
- S. D. Lawley, J. P. Keener. A new derivation of Robin boundary conditions through homogenization of a stochastically switching boundary. *SIAM Journal on Applied Dynamical Systems* 14 (2015) 1845–1867.
- R. D. Luce. *Individual Choice Behavior: A Theoretical Analysis*. Wiley, 1959.
- A. A. J. Marley, H. Colonius. The ‘horse race’ random utility model for choice probabilities and reaction times, and its competing risks interpretation. *Journal of Mathematical Psychology* 36 (1992) 1–20.
- D. McFadden. Conditional logit analysis of qualitative choice behavior. In P. Zarembka (ed.), *Frontiers in Econometrics*, Academic Press (1974) 105–142.
- D. McFadden, K. Train. Mixed MNL models for discrete response. *Journal of Applied Econometrics* 15 (2000) 447–470.
- G. A. Pavliotis, A. M. Stuart. *Multiscale Methods: Averaging and Homogenization*, Texts in Applied Mathematics 53, Springer, 2008. Chapter 9, Averaging for Markov chains.
- Z. Schuss, A. Singer, D. Holcman. The narrow escape problem for diffusion in cellular microdomains. *Proceedings of the National Academy of Sciences* 104 (2007) 16098–16103.

- A. Szabo, D. Shoup, S. H. Northrup, J. A. McCammon. Stochastically gated diffusion-influenced reactions. *Journal of Chemical Physics* 77 (1982) 4484–4493.
- A. Tsiatis. A nonidentifiability aspect of the problem of competing risks. *Proceedings of the National Academy of Sciences* 72 (1975) 20–22.
- A. Tversky. Elimination by aspects: a theory of choice. *Psychological Review* 79 (1972) 281–299.
- N. G. van Kampen. A cumulant expansion for stochastic linear differential equations I. *Physica* 74 (1974) 215–238.
- J. W. Vaupel, K. G. Manton, E. Stallard. The impact of heterogeneity in individual frailty on the dynamics of mortality. *Demography* 16 (1979) 439–454.
- J. I. Yellott Jr. The relationship between Luce’s choice axiom, Thurstone’s theory of comparative judgment, and the double exponential distribution. *Journal of Mathematical Psychology* 15 (1977) 109–144.
- G. G. Yin, Q. Zhang. *Continuous-Time Markov Chains and Applications: A Singular Perturbation Approach*. Springer, 1998.
- R. Zwanzig. Rate processes with dynamical disorder. *Accounts of Chemical Research* 23 (1990) 148–152.